AI IN THE COURTS · CASE 016

One researcher's database lists 2,185 court rulings on AI-hallucinated filings

A researcher's database lists 2,185 court decisions in which, he says, a court found or implied that a party relied on hallucinated content. More than half of its entries list a self-represented litigant, and about one entry in nine names a specific tool. Two state supreme courts show the range: in Nebraska the lawyer said he did not rely on AI and the court did not decide; in New Mexico the lawyer admitted using ChatGPT.

WHAT HAPPENED

- Damien Charlotin keeps a database of court and tribunal decisions which, he says, covers only cases where the court found or implied that a party relied on hallucinated content. As downloaded on 10 October 2026 it lists 2,185 decisions, the oldest from April 2023.
- More than half of the entries list a self-represented litigant as the party involved; 865 list a lawyer, and 32 list a judge.
- Only 234 entries name a specific tool. Most say the tool is "Implied" or "Unidentified", and the database itself warns that naming a tool does not necessarily mean that tool was responsible.
- In Nebraska, a lawyer whose brief cited a case that does not exist said he did not rely on AI. The state's Supreme Court did not decide whether he had: it struck the brief, dismissed the appeal and referred him for discipline, saying its analysis was the same either way.
- In New Mexico, the state's Supreme Court records that a lawyer admitted using ChatGPT for a brief with invented witnesses and testimony. The court held him in contempt, ordered him to pay $5,000 and referred him for discipline.

Grade B Well supported

KEY NUMBERS

2,185
court and tribunal decisions the database lists, April 2023 to October 2026

1,259
entries that list a self-represented litigant, against 865 that list a lawyer

About 1 in 9
entries that name a specific tool (234 of 2,185)

164
entries that record a professional sanction; 380 record a monetary penalty

WHY IT MATTERS

More than half of the decisions in this database list a self-represented litigant, and only about one in nine names a specific tool. The library's reading: the risk shown here is less any one product than unchecked work reaching a court, and the Nebraska Supreme Court held that its existing rules were enough to deal with it.

STILL UNKNOWN

- Whether AI produced the fictitious citations in the Nebraska brief. The court did not decide, and the lawyer denied it.
- Which tool, if any, lies behind most entries: 1,925 of them say only "Implied" or "Unidentified".
- How many fake citations or uses of AI in court filings the database does not capture. The compiler says it does not track that wider universe.

Read the full record, with every source and its quotes: https://danielwipert.github.io/ai.case.study/cases/2026/aai-2026-016-fabricated-legal-authority-court-sanctions/

Everything above except "Why it matters" is the case's checked summary text; "Why it matters" is the library's reading, checked against the same record. Checked 11 oct 2026. Made 25 Oct 2026.